Epileptic Seizure Detection from EEG Signals with Autoencoded Features

Tuan Nguyen

March 28, 2018

1 Definition

1.1 Project Overview

The domain of interest of this project belongs to physiological data such as electroencephalography (EEG), magnetoencephalography (MEG), electrocardiography (ECG) and the recorded signals from wearable devices. This project focuses on the EEG signals, which capture activities of brain neurons during a period of time. Different kinds of EEG data has been recorded from humans, for instance from those at rest, sleep [11], during some periods of specific cognitive activities, or from patients with diseases such as Alzheimer's, Parkinson's, depression and epileptic seisures ([1] and references thereof).

This project studies the classification problem on an EEG data set recorded from healthy volunteers and patients having epileptic seizures [1]; solutions for this problem could be used to assist with detecting patients with the disease from their brain activity signals. Previous work on this data set focused mainly on extracting hand-crafted features to be used for classification. As examples, Nigam and Graupe [12] extracted the spike amplitudes and frequency of the signals, feeding them into a neural network for classification; Guler et al. [6] applied wavelet transformation to the signals to extract features, and classification was performed using a neuro-fuzzy system; Kannathal et al. [9] extracted entropy-based features for classification. On another data set, Wulsin et al. [15] learned features of second-long segments of EEG signals using Deep Belief Networks [7] and classified them into

"clinically significant" EEG classes. To the best of my knowledge, there was no previous work reporting the results of using autoencoders for the data set that this project is interested in. The purpose of this work, therefore, is to explore the use of traditional and denoising autoencoders ([2], [14]) in learning the features representing the EEG signals and then use them to classify unseen brain activity signals into different classes of patients. For this classification purpose, the current implementation was limited to the use of multi-layer fully-connected neural networks where their hidden layers were used to fine-tune the feature learned by the autoencoders, and a softmax layer served as the ouput layer discriminating input signals into classes. To the end, I present extensive experiments that show the great benefits of using denoising autoencoders in detecting epileptic seizure from EEG signals.

1.2 Problem Statement

This project aims to classify 1-second long EEG segments into one of the three classes: (1) healthy volunteers, (2) patients with epileptic seizures disease during seizure-free periods, and (3) patients with the disease during active seizure periods. It can be formally defined as follows.

- **Input:** training set $X = \langle \mathbf{x}_1, \mathbf{x}_2, ..., \mathbf{x}_N \rangle$ of N samples, where $\mathbf{x}_i$ is a vector of M electrical voltage values recorded in one second by an electrode at a specific point and time; target vector $y = \langle y_1, y_2, ..., y_N \rangle$ where $y_i \in \{1, 2, 3\}$ is the type of volunteers which the sample $\mathbf{x}_i$ belong to. In the data set [1], the targets $1, 2, 3$ respectively correspond to sets of volunteers $A + B$, $C + D$ and E.

- **Output:** a hypothesis h classifying a sample of M electrical voltage values into one of the volunteer classes $\{1, 2, 3\}$.

The formation of the training set X and target y is presented in more details in Section 2.1. This project explores the benefits of using autoencoders in learning a representation $f(\mathbf{x}_i)$ of signal segments $\mathbf{x}_i \in X$ that can be helpful for our classification task at hand; the two types of autoencoders considered in this work are traditional [2] and denoising autoencoder [14]. The trained parameters of these autoencoders are then used to fine tune multi-layer feed-forward neural network for classification (see Section 2.2).

1.3 Metrics

The purpose of a hypothesis returned for the problem stated above is to classify unseen brain activity segments, correctly as frequently as possible, into one of the volunteer classes $\{1, 2, 3\}$. The *accuracy* is the most important measure for this purpose. Given a hypothesis h and a set of examples T, let $C_k(h, T)$ be the number of examples in T classified into correct class k ($k \in \{1, 2, 3\}$) by h. The accuracy of h on the set T is formally defined as follows:

$$Acc(h, T) = \frac{\sum_{k \in \{1,2,3\}} C_k(h, T)}{|T|}.$$

The accuracy measure is used on both the training set to update parameters of a model under construction and the validation set to approximate the performance of the model on future unseen data set. Candidate models are ranked based on their accuracy on an independent test set.

2 Analysis

2.1 Data Exploration and visualization

The data set used in this study is provided with the work by Andrzejak et al. [1], recording brain electrical activity of five sets of human volunteers:

- Set A: healthy volunteers in relaxed state with eyes open.

- Set B: healthy volunteers in relaxed state with eyes closed.

- Set D: epilepsy patients during seizure-free periods; the electrodes were implanted to record brain activity from within the epileptogenic zone.

- Set C: epilepsy patients during seizure-free periods; the electrodes were implanted from the opposite hemisphere of the brain (compared to those in Set D).

- Set E: same patients with sets C and D but activity were recorded from all electrodes during active seizures.

Class	Min	Max	Mean	Stddev	25%	50%	75%
A	-288	294	-6.26	48.33	-38	-6	25
B	-424	360	-12.51	70.68	-57	-12	32
C	-412	623	-8.88	59.38	-43	-7	26
D	-1147	2047	-6.20	90.34	-43	-6	29
E	-1885	2047	-4.74	341.15	-175	-5	177

Table 1: Statistics of the data set per class.

For each set of volunteers above, the data set contains 100 single-channel EEG segments of 4096 signal micro-Voltage values in time domain of 23.6-sec duration. For this project each of these segments is divided into smaller segments of 1-second durations, which are considered stationary for EEG data [12]. These 1-second durations together define the training set X of 178 dimensions, and the three sets of volunteers $A + B$, $C + D$ and E define the target vector y in the problem statement above.

Figure 1 shows four randomly chosen 23.6-second segments of each class from which we can only observe very slightly differences in patterns of the signals in each class. However, by looking at the shape of each class more closely it is interesting to notice some very distinguished segments that belong to epileptic seizure patients in both seizure-free or seizure-active periods (i.e., class C, D and E), as shown in Figure 2. We can observe that some signals that belong to patients during their seizure-active periods (i.e., class E) are very dense, whereas classes C and D contain some signals that are more sparse but also include several high peaks within one second.

Figure 3 shows the chart of common statistics of signal values within each class, and the detailed numbers are shown in Table 1. The data points of class E, the category of patients in their seizure-active periods, have the smallest average value (-4.74) but spread the most among the classes with standard deviation 341.15 (more than the total standard deviation of all the other classes). The four classes A, B, C and D appear to have similar percentiles of 25%, 50% and 75%; however, data points in class D appear to spread wider than those in the other three classes. The ranges of minimal and maximal values are perhaps the most visible distinction between the two sets $A \cup B$ and $C \cup D$.
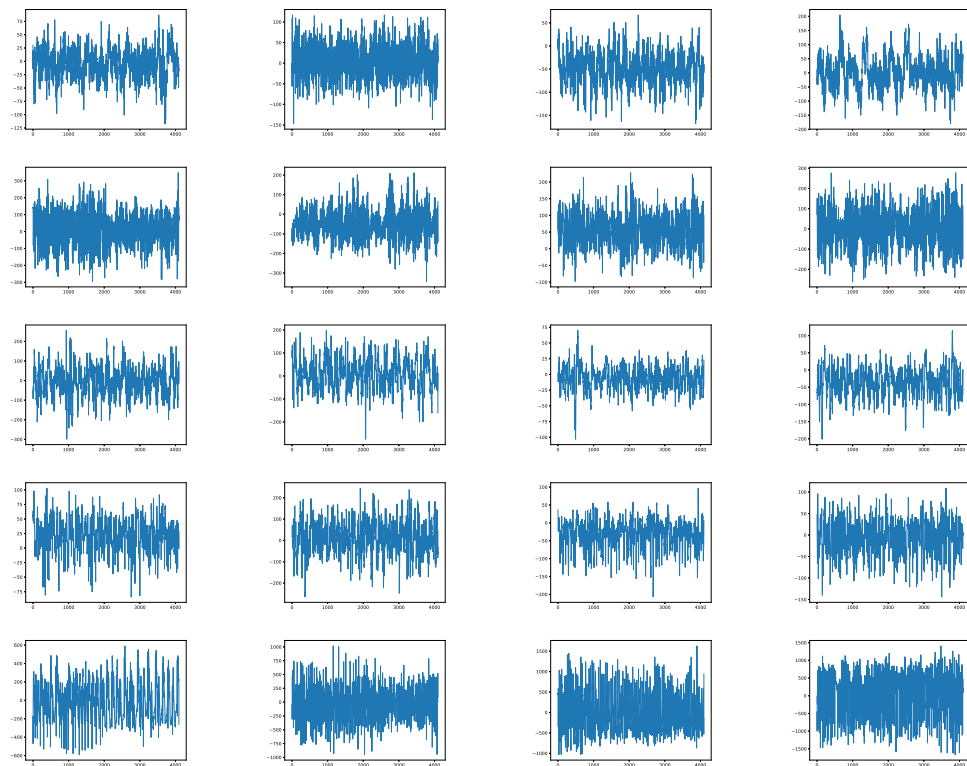
4

Figure 1: Images of brain activity signals from patients of classes A (1st row), B(2nd row), C(3rd row), D(4th row) and E(5th row).
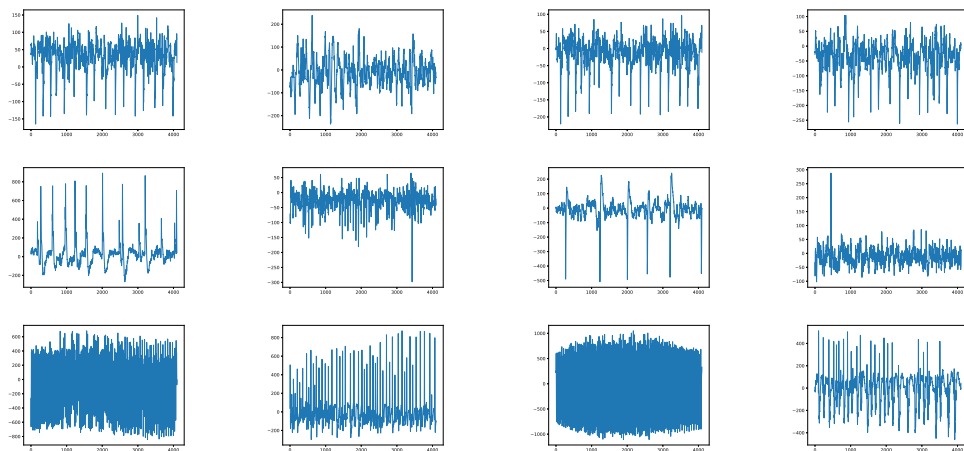
Figure 2: Some distinguished signals from patients of classes C (1st row), D(2nd row), E(3rd row).
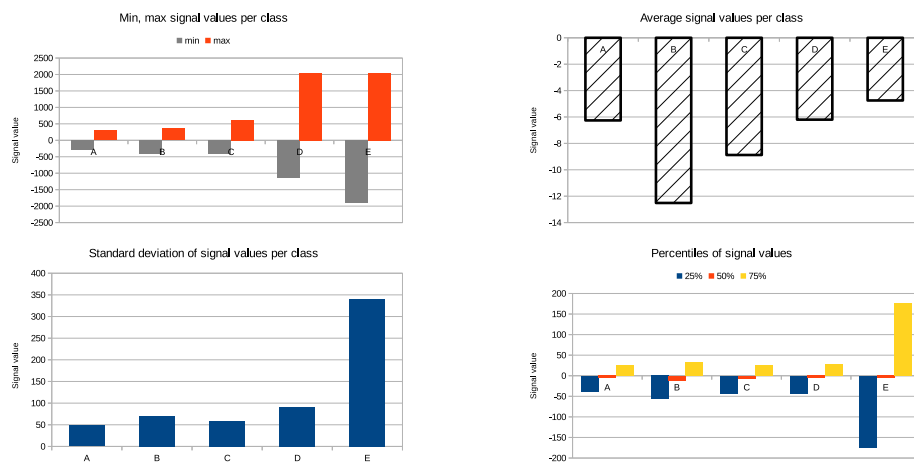

Figure 3: Some statistics of the data set per class: min, max (top left), mean (top right), standard deviation (bottom left) and percentiles (bottom right).

2.2 Algorithms and Techniques

As mentioned earlier, much work has been done on this data set that focused on extracting hand-crafted features to be used for classifying the segments into volunteer target classes ([12], [6], [9]). Since the ability to automatically learn features and representations is a key benefit of machine learning, a reasonable question is whether one could automatically learn features from the EEG signals in this domain for classification purposes, and if so then which algorithm could be used effectively.

Deep architectures of neural networks have been successful recently in many machine learning applications, and part of their success is due to their ability to learn high-level features. On the other hand, training deep networks is costly, and one way to tackle this challenge is to initialize their parameters with values obtained from pretraining some types of networks. Two types of such networks have been widely used: restricted Boltzmann machines [7] and autoencoders[2]. For EEG signal domains, previous work such as Wulsin et al. [15] pretrained restricted Boltzmann machines for feature learning and classification, and although Bengio et al. [2] showed that autoencoders worked equally well as restricted Boltzmann machines, there were no previous work, to the best of my knowledge, which used autoencoders for learning representations and classification in the epileptic seisure domain [1] of interest.

The proposed approach therefore is to train autoencoders for extracting representations of the input signals, and then ultilize the pretrained autoencoders to form a deep neural network for the classification task at hand. The architecture of multi-layer feed-forward neural networks was chosen as (1) it can nicely incorporate the encoding parts of the pretrained autoencoders, and (2) it allows for *fine-tuning* the pretrained parameters for classifying the specific target classes, a technique that usually help improve the performance of deep neural networks.

2.2.1 Training autoencoders

A *traditional autoencoder* ([2], [14]) is a fully connected feed-forward neural network which consists of an input layer, one hidden layer and an output layer such that the input and output layers have the same number of neurons. The computation performed by the hidden layer, specified by a weight matrix $\mathbf{W}$, bias vector $\mathbf{b}$ and nonlinear activation function s, acts as an *encoder* that

maps an input vector $\mathbf{x}$ to a representation $f_\theta(\mathbf{x})$ defined as follows:

$$f_\theta(\mathbf{x}) = s(\mathbf{W}\mathbf{x} + \mathbf{b}).$$

The parameter set $\theta = (\mathbf{W}, \mathbf{b})$ is part of what to be adjusted. The encoding $f_\theta(\mathbf{x})$ is then mapped back to vector $\mathbf{z}$ in the same space with $\mathbf{x}$ through the computation of the output layer, representing the *decoder* part of the autoencoder:

$$\mathbf{z} = g_{\theta'}(f_\theta(\mathbf{x})),$$

where $g_{\theta'}$ can be either a linear or a non-linear function whose parameter θ' is the other part of what to be learned. With real-valued input signals $\mathbf{x}$, the output $\mathbf{z}$ in general is interpreted not as the exact reconstruction of $\mathbf{x}$ but the mean of a distribution $P(\mathbf{x}|\mathbf{z})$. Squared error function between $\mathbf{x}$ and $\mathbf{z}$, and thus the objective function

$$L(X) = \frac{1}{N}\sum_{i=1}^{N}(\mathbf{x_i} - \mathbf{z_i})^2,$$

is to be optimized during the training of the autoencoders. As suggested in [14], affine mapping is to be used for this type of real-valued inputs and error function; in other words, $g_{\theta'}(\mathbf{y}) = \mathbf{W}'\mathbf{y} + \mathbf{b}'$ and $\theta' = (\mathbf{W}', \mathbf{b}')$.

In order for autoencoders not to produce trivial encoding and decoding functions (e.g., to avoid the identity function from input $\mathbf{x}$ to itself), certain types of constraints need to be imposed. In traditional autoencoder, a constraint could be that the number of neurons of the hidden layer is smaller than the dimension of the inputs $\mathbf{x}$. In so-called *denoising autoencoders* [14], the inputs are made "dirty" before being fed into the encoder—by attempting to reconstruct the original input from its corrupted version, denoising autoencoders might learn representations that are more essential to the inputs. In this work, the input signals are made corrupted using one of the following two ways:

- Adding into the inputs random noises generated by a Gaussian distribution with zero mean and σ standard deviation $N(0, \sigma^2)$.

- Disable a random number of neurons at the input layer of the autoencoders, which can be implemented using dropout techniques [13] with dropout rate denoted by ρ $(0 < \rho < 1)$.

2.2.2 Stacking up autoencoders

The idea of reconstructing the inputs using autoencoders can be applied to map the orginal input into a highly non-linear representations. Specifically, the process starts by training the first autoencoder, defined by encoding and decoding functions $f_{\theta_1}^{(1)}$ and $g_{\theta_1'}^{(1)}$, to reconstruct input vector $\mathbf{x}$. The representation of $\mathbf{x}$ is then obtained by applying the learned encoding function to $\mathbf{x}$, i.e. $f_{\theta_1}^{(1)}(\mathbf{x})$. For each index $i \geq 1$, the representation $f_{\theta_i}^{(i)}(\mathbf{x})$ produced from training the ith autoencoder becomes the input to the next $(i + 1)$-th autoencoder, which is then trained to adjust its parameter θ_{i+1} and θ_{i+1}' of the encoding and decoding functions $f_{\theta_{i+1}}^{(i+1)}$ and $g_{\theta_{i+1}'}^{(i+1)}$, respectively. From the resulting sequence of trained encoders $\langle f_{\theta_1}^{(1)}, f_{\theta_2}^{(2)}, ..., f_{\theta_k}^{(k)} \rangle$, the original input $\mathbf{x}$ can then be mapped into the representation computed as $f_{\theta_k}^{(k)}(f_{\theta_{k-1}}^{(k-1)}(...(f_{\theta_1}^{(1)}(\mathbf{x}))))$.

More importantly, a multi-layer feed-forward neural network can be built based on the sequence of trained autoencoders for our classification task. The network consists of the following components:

- an input layer of M neurons, where M is the number of dimension of the input vectors, that accepts input signals $\mathbf{x}_i \in X$;

- a sequence of k hidden layers where the ith layer $(i \geq 1)$ is the hidden layer of the i-th pretrained autoencoder, i.e. *the weights and biases are initialized with the pretrained values*;

- a softmax output layer with three neurons for the three target classes *whose weights and biases are randomly initialized.*

The above network can be used in two ways. In the first approach, the weights and biases of all the hidden layers are freezed, and the parameters of the output layer are to be trained subject to the reconstruction error defined on the training set X; in other words, the representations $f_{\theta_k}^{(k)}(f_{\theta_{k-1}}^{(k-1)}(...(f_{\theta_1}^{(1)}(\mathbf{x}))))$ are used directly as the inputs to the output layer whose parameters are the only ones to be adjusted by the network during learning. The second approach, which results in better accuracy, is fine-tuning the entire network's parameters for the classification task. By starting from the pretrained weights and biases in the hidden layers, we can train an instance of the above networks with competitive accuracy performance, compared to other methods solving this problem using hand-crafted features.

2.3   Benchmark

The main purpose of this work is to explore the benefit of using pretrained autoencoders to train a multi-layer feed-forward neural network for the classification task at hand. Secion 4 presents several results between different design choices and network configurations. For external comparisons, previous work has been developed for the classification problem on the same data set [1]; the main contribution of most of these work, however, were on designing various hand-crafted features for the EEG signals. They might also be different from each other in the target classes of interest. As examples, Nigam and Graupe [12] used two features based on the spike information of the signals for further learning and detecting patients in set E (thus, binary classification problem). Kabir et al. [8] extracted statistical features based on "optimum allocation technique" to be used with logistic model trees for classifying volunteers into the three classes that are of interest of this project. Guler et al. [6] employed Lyapunov exponents based features for classification using recurrent neural networks where the target classes are three sets A, D and E. On another data set, Wulsin et al. [15] learned features of second-long segments of EEG signals using Deep Belief Networks [7] and classified them into "clinically significant" EEG classes. To the best of my knowledge, there was no previous work reporting the results of using autoencoders for the data set that this project is interested in.

The proposed method by Kabir et al. [8] on the data set of interest in this project outperforms many the state-of-the-art approaches using the same data set; among many other measures, the total accuracy of their approach was 95.3%. Since its target classes are the same with those of interest in this project, and their work represents a class of work using hand-crafted features, it is reasonable to use the performance of their work as the baseline for the approach to be explored in this project. The hypothesis is that by using autoencoders, it is possible to learn a set of features automatically for classifying EEG signals with comparable results.

3   Methodology

3.1   Data Preprocessing

The data set [1] contains 100 channels (vectors of micro-voltage values) for each of the five sets of volunteers. Each channel is a vector of 4097 values that

are EEG signals recorded in 23.6 seconds. The data set is then preprocessed as follows.

- Each channel is divided into 23 one-second segments with equal length, each contains 178 values. These segments are marked with the target value 0, 1 or 2 for the classes $A + B$, $C + D$ and E of the original channel.

- The resulting set of segments are splitted into three separate sets for training (60%), validation (20%) and testing (20%) (using Stratified-ShuffleSplit function from Scikit-Learn).

- The resulting set of segments in the training set are then normalized so that the values of EEG signals are in between $[-1, 1]$. The internal state of the scaler (which contains for example the min and max signal values of the training set) is used to scale the other two sets into the same interval.

3.2   Implementation

The approach proposed in this project was implemented in Python 3 using TensorFlow library, and its source code could be found at the following link:

$$https://github.com/natuan/ml\_nano\_capstone.git.$$

In the following, I discuss at high level the key components of the implementation of autoencoders, the derived feed-forward neural network, and the elements that support for their training and testing.

3.2.1   Autoencoders

The autoencoders were implemented in the class `UnitAutoencoder` inside folder *src/StackedAutoencoders.py*. Figure 4 shows the conceptual image of the two types of autoencoders, optionally with dropout applied to the hidden layers for dealing with overfitting. In TensorFlow, the key elements of the autoencoders can be implemented roughly as follows.

Input tensors: the input vectors of size `n_inputs` are placed into a "place holder"[1]:

[1] Here `tf` stands for the imported module tensorflow.

11

Figure 4: Traditional (left) and denoising (right) autoencoders with an input layer, a hidden layer and an output layer. Dropout technique is applied to hidden layers in the two types of autoencoder for dealing with overfitting, and to the input layer also for corrupting the input vectors.

```
X = tf.placeholder(tf.float32,
                   shape=[None, n_inputs],
                   name='X'),
```

whose value will be assigned at run-time. To facilitate the use of dropout, which takes effect only during training, one needs also define a boolean tensor as another input of the network indicating the training/test mode:

```
training = tf.placeholder_with_default(False,
                                       shape=(),
                                       name='training').
```

Input layer: the input layer of the traditional autoencoders is specified simply with the tensor X. When dropout is used to corrupt the inputs for denoising autoencoders, the dropout layer construct in TensorFlow, `tf.layers.dropout`, can be used on top of the original input if the `training` is set to be `True`:

```
X_noisy = tf.layers.dropout(X, input_dropout_rate,
                            training=training),
```

where the second parameter specifies the dropout rate. Alternatively, Gaussian noises, implemented using `tf.random_normal`, can be used to corrupt the inputs if the network is in training mode:

```
X_noisy = tf.cond(training,
                  lambda: X + tf.random_normal(tf.shape(self.X),
                                               stddev = noise_stddev),
                  lambda: X).
```

12

Hidden layer: The construct `tf.layers.dense` can be used to construct the dense hidden layer of `n_neurons` neurons right after `X_noisy` as follows:

```
dense_hidden = tf.layers.dense(X_noisy, n_neurons,
                               activation=hidden_activation,
                               ...),
```

where `hidden_activation` specifies the activation function used for the neurons. (Unused arguments were omitted to simplify discussion.) A dropout layer could then be put on top of the hidden layer for dealing with the overfitting issue:

```
hidden = tf.layers.dropout(dense_hidden,
                           hidden_dropout_rate,
                           training=training),
```

where `hidden_dropout_rate` defines the dropout rate for the hidden layer.

Output layer: the output layer of `n_inputs`, the same size with the input layer, can be implemented using `tf.layers.dense`, right after the layer `hidden` above:

```
outputs = tf.layers.dense(hidden,
                          n_inputs,
                          activation=output_activation,
                          ...).
```

The `output_activation` can be used to customize the activation function for the output layer; as mentioned earlier an affine mapping should be used given the real-valued inputs and the squared reconstruction error function.

Reconstruction error: in order to extract representations of the inputs, the autoencoder is trained to minimize the difference between its outputs and the original inputs. The reconstruction error is defined using the squared-error loss:

```
reconstruction_loss = tf.reduce_mean(tf.square(outputs - X)).
```

Training operator: To train the network, a training operator first needs to be defined to optimize the above reconstruction loss:

```
training_op = optimizer.minimize(reconstruction_loss).
```

13

Each step of the training is performed by running the training operator within a TensorFlow session `sess`, feeding the network with a current batch `X_batch` and the `True` value for input tensors `X` and `training`, respectively:

```
sess.run(training_op, feed_dict={X: X_batch,
                                 training: True}).
```

3.2.2 Building a feed-forward neural network from pretrained autoencoders

The feed-forward network is largely built based upon the encoding parts of the pretrained autoencoders, in particular for the architecture and parameters of its hidden layers. While the autoencoders are trained in an unsupervised fashion (their inputs are the expected outputs), the network is designed for supervised learning with predefined target classes. The network's tensor inputs, therefore, include the class tensor `y` (in addition to the tensor representing the input vectors `X` and the training/testing mode `training`):

```
y = tf.placeholder(tf.int64, shape=(None), name="y")
```

Softmax output layer

A copy of the original data set [1] can be found in folder **input**, and the training/validation/test sets can be found in folder **cache**. The main parts of the source code, located inside folder **src**,are as follows:

- **DataSet.py**: class `DataSet` manages the loading and processing of the data set. Its tasks include dividing the original into smaller segments and assigning them into appropriate target classes. The function `create_data_set_class_AB_CD_E()` could be invoked to re-generate the training, validation and test sets with customized a split ratio and random seed. (The `root_dir` local variable in this function needs to be changed to refer to the project folder.)

- **StackedAutoencoders.py**: contains the classes of autoencoders, `UnitAutoencoder`, and multi-layer feed-forward neural networks, `StackedAutoencoders`, built upon them. It also includes a builder class `StackBuilder` that is used in `Main.py` to construct and train individual autoencoders before using them to build the stacks (i.e., the multi-layer feedforward neural networks).

14

- **Main.py**: the main file where experiments were configured to run. The `_main_` function contains code to define autoencoders and the corresponding neural networks, as well as how to fine-tune and test the network on the traininig, validation and testing sets. Before running this file (as `>> python3 Main.py`), the `root_dir` local variable needs to point to the project folder.

3.3   Refinement

Two refinements were made: replacing sigmoid with ELU, discover bug regarding dropout (by looking at the learning curve)
I spent much time training networks with sigmoid activation functions for hidden layers without much success (various values for the standard deviation of the Gaussian noise distribution and dropout rates were also used). Although the reconstruction errors of autoencoders were small enough, the accuracy performance of the corresponding neural networks was never more than 60% even on the training set (with networks having 2 hidden layers, each with 200 neurons). I suspected that the shape of the sigmoid functions at some point during the training made the gradients extremely small (similar to observations reported in [5]), and therefore learning became very slow. I eventually settled with the ELU activation function [4] that appeared to work well for the classification task at hand.

4   Results

The following hyper-parameters and settings were used throughout various configurations whose results are reported in this section.

- ELU activation function [4] was used for hidden layers of the autoencoders and the corresponding multi-layer feed-forward neural networks. As mentioned above, it provided much better result compared to the sigmoid function.

- The Adam optimizer was used with the learning rate of 5e−6. Quick but incomplete experiments with higher learning rates result in networks that more quickly overfitted the training set.

- Batch size was chosen to be 64.

- Number of epochs was set to be 3×10^6, but early stopping was used when learning curves on the validation set started to go up, indicating that the corresponding model began to overfit the training set. (Many of the experiments below took about 2 days to complete.)

- Check points were defined as the batch update that was a multiple of $5 \times batchsize$, at which the accuracy performance was computed on the training and validation sets.

- Dropout technique was used consistently in all network configurations to deal with overfitting. As suggested by Srivastava et al. [13], common values 0.33 and 0.5 were used for the dropout rates of neurons respectively in the input layer (unless noises were used to corrupt the input) and the hidden layers.

- To reduce the number of configurations that need to explore, I restricted to consider only multi-layer feed-forward neural networks that have the same number of neurons across different hidden layers.

In the following discussion, the network configurations are named based on the first subsection in which they are described.

A. Randomly initialized vs. pretrained weights and biases: We first examine the benefit of using weights and biases pretrained in autoencoders to initialize a multi-layer neural networks. Figure 5 shows the learning curves for the training set (in red) and the validation set (in green) for a stack of two denoising autoencoders that were not pretrained. In other words, the weights and biases of the 2-hidden-layer fully connected network with 50 neurons per layer, called A_0, were initialized randomly. Dropout technique with the rate 0.33 (for input layer) and 0.5 (for hidden layers) are used to cope with overfitting. During the training of this network, the best model was saved at step $3,878,215$ of batch updates, after which the network started to overfit to the training data (i.e., the validation curve began to go up).

Figure 6 shows the learning curves for a similar stack as above but the two denoising autoencoders were pretrained to reconstruct their inputs, and the resulting weights and biases were used to initialize the stack before being fine-tuned. We could observe that the network did not overfit even after more than 18 millions batch updates, which is much better than A_0. The cross-entropy error of this network on the validation set can also be seen as

Figure 5: Cross-entropy error (Y-axis) on the training (red) and validation (green) sets of a 2-hidden-layer, 50 neurons each, fully connected network, called A_0, with randomly initialized weights and biases. Notice that the validation curve started to go up after about 3.8 millions of batch updates, indicating the overfitting of the network.

Network	Train set	Validation set	Test set
A_0	89.59%	87%	87.52%
A_1	93.17%	88.96%	90.34%

Table 2: Accuracy by networks A_0 and A_1 on training, validation and test sets.

much better than from the one with randomly initialized weights and biases: less than 0.3 vs. greater than 0.4. This by definition results in a difference between A_1 and A_0 in terms of prediction accuracy on the validation, and importantly, the *test* sets: 90.34% vs. 87.52%, respectively. The accuracy errors by the two networks on the various sets are shown in Table 2.

Although the network A_1 had a wider accuracy gap between its training and validation sets, it is important to stress that this gap did not appear to be widening, which was not the case for the network A_0. If the training were to be allowed on A_1, we would expect that its generalization capacity continues to be enhanced; the network A_0, however, overfitted early and would not have any benefits from further training.

B. Traditional vs. denoising autoencoders: We now compare two types of autoencoders, traditional and denoising, in terms of reconstructing orginal inputs and, more importantly, of learning representations for our classifica-

17

Figure 6: Cross-entropy error (Y-axis) on the training (red) and validation (green) sets of a 2-hidden-layer fully connected network, called A_1, whose weights and biases were initialized with pretrained values.

tion problem. The denoising autoencoders here were those used to pretrain weights/biases of A_1, which learned their unique features by reconstructing the inputs from their corrupted version. The traditional autoencoders had the same configuration with the other type (i.e., number of neurons, dropout rates for the hidden layer), except that they attempted to reconstruct the inputs from themself instead of their corrupted version. To test their representations in classifying the target signals, we constructed a 2-hidden-layer fully-connected neural network, called B_0, whose weights and biases were initialized from the two pretrained traditional autoencoders. No Gaussian noises or dropout was applied to the inputs of the network B_0.

Table 3 shows the reconstruction errors by two autoencoders of the above types used to construct the layers of networks A_1 and B_0. On both the training and validation sets, the traditional autoencoders which processed the inputs directly resulted in more accurate reconstructed signals, i.e. with smaller errors, than those produced by the denoising counterpart; and this is true for autoencoders corresponding to both two hidden layers of the networks.

Figure 7 shows the 522th segment and its reconstructed signals produced by networks B_0 and A_1. We can notice that at around the dimension 150th, the signal generated by the network B_0 shows more accurate reconstructed curve than that of A_1. This observation can be seen with other segments and their reconstructed signals, consistent with the comparison above between the networks in terms of their reconstructed errors.

The outperformance of traditional autoencoders in reconstructing the input signals, however, did not transform to its classification performance.

18

Autoencoder	Layer 1		Layer 2	
	Train	Validation	Train	Validation
Traditional	0.0015	0.0016	0.00051	0.00052
Denoise	0.0019	0.0021	0.0011	0.0011

Table 3: Reconstruction errors by traditional and denoise autoencoders used at layer 1 and 2 of networks B_0 and A_1, respectively.


Figure 7: Segment 522th and its reconstructed signals produced by networks B_0 (left) and A_1 (right).

Figure 8: Cross-entropy error (Y-axis) on the training (green) and validation (red) sets of a 2-layer fully connected network, called B_0, with inputs not being corrupted during training.

Figure 8 shows the cross-entropy errors by the network B_0: it was expressive enough and achieved 97.13% accuracy on the training set, however overfitted early after step $1,675,620$ of batch updates. The accuracy of the resulting best model on the validation and testing sets were 90.47% and 91.78%, respectively. (The corresponding statistics by A_1 was presented in the previous result section.) Both the about 7% performance gaps on between the training and the two independent sets, and the $> 1\%$ gap between the validation and test sets indicate that this model suffers from high variance. The network A_1, which used features extracted through reconstructing corrupted inputs, as shown in Figure 6, provided a comparable and more reliable accuracy performance on unseen data.

Perhaps one explanation for the above difference could be that since the two networks were the same except for the clean vs. the corrupted inputs that they processed, the noises introduced into the inputs helped the network A_1 learned features more essential to the original inputs; when used for classification task, these features resulted in much less overfitting model compared to those simply learned from an identity function of the clean inputs. Although it is hard for humans to make sense the features learned by the two networks to represent brain activity signals (in contrast to interpretable features in image domains), it is interesting to see what the neurons were "excited" about after the learning process. Figure 9 shows the weights of six randomly chosen neurons learned by the first traditional autoencoder used in network B_0: we could observe that the weights did not appear to learn any interesting patterns after reconstructing the original inputs.

20

Figure 9: Weights of six neurons, chosen randomly, learned by the first traditional autoencoder used in network B_0 after reconstructing the original (uncorrupted) inputs.

Network	Train set	Validation set	Test set
B_1	97.3%	89.61%	90.22%
B_2	94.28%	89.99%	90.69%

Table 4: Accuracy by networks with 2 layers of 75 neurons B_1 and B_2 on training, validation and test sets. B_1 reconstructed the original signals from themselves, and B_2 did so from the corrupted inputs.

For the same set of neurons, Figure 10 shows their weights learned by the denoising autoencoder used in the first hidden layer of network A_1 after reconstructing the inputs from their corrupted version. As opposed to those in Figure 9, these weights appear to reflect some specific patterns that the neurons were adapted to. Similar contrasts could also be observed for autoencoders at the second layers of the networks B_0 and A_1.

Comparison on networks of 75-neuron layers: Table 4 shows the accuracy performances of networks B_1 and B_2 with 2-hidden-layer fully connected networks, each having 75 neurons. They differed in exactly the same way B_0 and A_1 were distinguished: B_1 reconstructed the original signals from clean inputs, whereas B_2 did that from the corrupted inputs. We can again observe from the accuracy gaps that by reconstructing the inputs from their corrupted version, the denoising autoencoders extracted features that helped the network classifier learned a model that was more reliable in predicting the target classes.

C. Gaussian noisy inputs vs. dropout inputs:
I found that using Gaussian noises to corrupt the inputs of autoencoders seemed to work equally well compared to using the dropout technique. Figure 11 shows the learning curves on the training and validation sets produced by a network, called C_0, that is similar to A_1 except for inputs being corrupted using Gaussian noises of 0.1 standard deviation (instead of dropout with rate 0.33 used by A_1 and its autoencoders.) The best model by C_0 was saved at step $17,575,820$ of batch updates, after which the performance on the validation set started to decrease indicating that the model began to overfit to the training data. The accuracy performance of C_0, compared to A_1, is summarized in Table 5.

Figure 10: Weights of the corresponding six neurons in Figure 9 learned by the first denoising autoencoder used in network A_1 after reconstructing the corrupted inputs.

Network	Train set	Validation set	Test set
C_0	93.84%	90.43%	90.52%
A_1	93.17%	88.96%	90.34%

Table 5: Accuracy by networks C_0 and A_1 representing the effects of using Gaussian noises and dropout to corrupt the inputs for learning representations.

Figure 11: Cross-entropy error (Y-axis) on the training (red) and validation (green) sets of a network, called C_0, that is similar to A_1 except for the inputs being corrupted using Gaussian noises instead of dropout.

Network	#Neurons/layer	Train set	Validation set	Test set
A_1	50	93.17%	88.96%	90.34%
B_2	75	94.27%	89.99%	90.69%
D_0	100	95.68%	90.56%	91.52%
D_1	250	99.42%	90.95%	91.95%

Table 6: Accuracy by networks of two hidden layers with different number of neurons per layer.

D. Denoising autoencoders with dropout: number of neurons and layers

Number of hidden neurons: Table 6 shows the accuracy performance on the three sets of 2-hidden-layer networks with different number of hidden neurons. We can see that increasing the size of the hidden layers from 50 to 75 and to 100 improved the accuracy up to 1%; unfortunately the performance gaps between training and validation/test sets also got wider, increasing the variances of the bigger networks. With 250 hidden neurons (network D_1), the network achieved 99.42% accuracy on the training set but did not improve much on the independent test set (91.95%). As shown in Figure 12, this network overfitted to the training set at step $6,073,855$ of batch updates. For that many hidden neurons, perhaps a more agressive dropout would be needed to cope with overfitting.

24

Figure 12: Cross-entropy error (Y-axis) on the training (red) and validation (green) sets of 2-hidden-layer network D_1 which has 250 hidden neurons per layer.

Network	#Hidden layers	Train set	Validation set	Test set
D_2	1	93.34%	89.26%	89.43%
A_1	2	93.17%	88.96%	90.34%
D_3	3	92.98%	88.56%	89.69%

Table 7: Accuracy by networks of 1, 2 and 3 hidden layers with 50 hidden neurons per layer.

Number of hidden layers: As shown above for 2-hidden-layer networks, increasing number of hidden neurons could help with improving the accuracy performance (though not much unless more agressive ways were used to deal with overfitting). The number of hidden layers is also another hyperparamter that determines the expressiveness of a neural network. With our small training set of about 6900 examples, we might not need much deeper neural network. Figure 13 and 14 shows the learning curves by networks of 1 and 3 hidden layers, each of which has 50 hidden neurons, and Table 7 shows the accuracy performance on the training, validation and testing sets of these networks compared to their counterpart that has 2 hidden layers. We can observe that their performance are quite close with A_1, which has 2 hidden layers, being slightly better than the others on an independent test set.

External comparison: As mentioned earlier, the work by Kabir et al. [8] achieved an accuracy of 95.3% on an independent test set. The best performance I could get from training various network configurations is about

Figure 13: Cross-entropy error (Y-axis) on the training (red) and validation (green) sets of a network, called D_2, that is similar to A_1 but has 1 hidden layer.


Figure 14: Cross-entropy error (Y-axis) on the training (red) and validation (green) sets of a network, called D_3, that is similar to A_1 but has 3 hidden layers.

90.34% to 91.52% (with 2-hidden layer networks of 50 and 75 neurons per hidden layer), which in my opinion is very encouraging for an approach that learns features automatically. One important difference between my work and the one by Kabir et al. [8] is that their approach required the original 23.6-second channels to be divided into 5.9-second segments, more than 5 times longer than those used in the present approach. Their method, which was based on computing special statistics from those 5.9-second segments, might need such long segments in order to gather enough information for classification.

5   Conclusion

This work presented and implemented an approach to classifying brain activitiy signals [1] into classes related to epileptic seizure disease. The proposed method resulted in neural network that achieved accuracy of about 90.34% to 91.52% on an independent test set. Although this accuracy result is lower than the work by Kabir et al. [8] on the same target classes, this work used much smaller 1-second segments, making it more sensitive to be used in potential applications such as medical wearable devices.

Due to my limited time, there are still items that I proposed but have not completed, and are worth further investigation. Future work include:

- consider using other types of autoencoders such as sparse autoencoders [3] and variational autoencoders [10] for this data set;

- consider using other classifiers such as Decision Trees, K-Kearest Neighbors, Support Vector Machines and Boosting with the representations learned by the sequence of autoencoders;

- compare different approaches in terms of other measures such as F1 score.

References

[1] R. G. Andrzejak, K. Lehnertz, F. Mormann, C. Rieke, P. David, and C. E. Elger. Indications of nonlinear deterministic and finite-dimensional structures in time series of brain electrical activity: Dependence on recording region and brain state. *Physical Review E*, 64(6):061907, 2001.

[2] Y. Bengio, P. Lamblin, D. Popovici, and H. Larochelle. Greedy layerwise training of deep networks. In *Advances in neural information processing systems*, pages 153–160, 2007.

[3] Y.-l. Boureau, Y. L. Cun, et al. Sparse feature learning for deep belief networks. In *Advances in neural information processing systems*, pages 1185–1192, 2008.

[4] D.-A. Clevert, T. Unterthiner, and S. Hochreiter. Fast and accurate deep network learning by exponential linear units (elus). *arXiv preprint arXiv:1511.07289*, 2015.

[5] X. Glorot and Y. Bengio. Understanding the difficulty of training deep feedforward neural networks. In *Proceedings of the thirteenth international conference on artificial intelligence and statistics*, pages 249–256, 2010.

[6] N. F. Güler, E. D. Übeyli, and I. Güler. Recurrent neural networks employing lyapunov exponents for eeg signals classification. *Expert systems with applications*, 29(3):506–514, 2005.

[7] G. E. Hinton and R. R. Salakhutdinov. Reducing the dimensionality of data with neural networks. *science*, 313(5786):504–507, 2006.

[8] E. Kabir, Siuly, and Y. Zhang. Epileptic seizure detection from eeg signals using logistic model trees. *Brain informatics*, 3(2):93–100, 2016.

[9] N. Kannathal, M. L. Choo, U. R. Acharya, and P. Sadasivan. Entropies for detection of epilepsy in eeg. *Computer methods and programs in biomedicine*, 80(3):187–194, 2005.

[10] D. P. Kingma and M. Welling. Auto-encoding variational bayes. *arXiv preprint arXiv:1312.6114*, 2013.

[11] M. Längkvist, L. Karlsson, and A. Loutfi. Sleep stage classification using unsupervised feature learning. *Advances in Artificial Neural Systems*, 2012:5, 2012.

[12] V. P. Nigam and D. Graupe. A neural-network-based detection of epilepsy. *Neurological Research*, 26(1):55–60, 2004.

[13] N. Srivastava, G. E. Hinton, A. Krizhevsky, I. Sutskever, and R. Salakhutdinov. Dropout: a simple way to prevent neural networks from overfitting. *Journal of machine learning research*, 15(1):1929–1958, 2014.

[14] P. Vincent, H. Larochelle, I. Lajoie, Y. Bengio, and P.-A. Manzagol. Stacked denoising autoencoders: Learning useful representations in a deep network with a local denoising criterion. *Journal of Machine Learning Research*, 11(Dec):3371–3408, 2010.

[15] D. Wulsin, J. Gupta, R. Mani, J. Blanco, and B. Litt. Modeling electroencephalography waveforms with semi-supervised deep belief nets: fast classification and anomaly measurement. *Journal of neural engineering*, 8(3):036015, 2011.